

A Detailed Exposition of the Bellare-Micali 1-out-of-2 Oblivious Transfer Protocol

Debrup Chatterjee

January 30, 2025

Abstract

Oblivious Transfer (OT) is a fundamental cryptographic building block with numerous applications in secure computation and privacy-preserving protocols. The Bellare-Micali 1-out-of-2 OT protocol is a classic and elegant solution under the Decisional Diffie-Hellman (DDH) assumption, leveraging ElGamal-like exponentiation and hashing. In this document, we present a detailed derivation of the protocol, step-by-step mathematical exposition, and an analysis of its security.

Contents

1	Introduction	2
1.1	Background and Motivation	2
1.2	Applications and Significance	2
1.3	Historical Development	3
1.4	Notation and Assumptions	3
1.5	Security Goals	4
2	Protocol Definition	4
2.1	High-Level Idea	4
2.2	Protocol Steps	5
2.3	Pseudocode	6
3	Correctness Proof	6

4	Security Analysis	8
4.1	Sender's Security	8
4.2	Receiver's Security	9
4.3	Extensions and the Malicious Setting	11
5	Conclusion	11
A	Mathematical Details and Implementation Considerations	12
A.1	Properties of XOR Operation	12
A.2	Group Selection and Security Implications	12
A.3	Parameter Size Recommendations	13
A.4	Efficiency Analysis	13
A.5	Implementation Considerations	13
A.6	Performance Optimizations	14
B	Security Analysis Extensions	14
B.1	Malicious Adversary Considerations	14
B.2	Advanced Security Properties	14

1 Introduction

1.1 Background and Motivation

Oblivious Transfer (*OT*) stands as one of the most fundamental primitives in modern cryptography, serving as a building block for numerous secure computation protocols. First introduced by Rabin in 1981, and later refined by Even, Goldreich, and Lempel, OT has evolved from a theoretical concept to a practical tool essential for privacy-preserving protocols.

In its most basic form, a 1-out-of-2 OT protocol involves two parties: a *sender* who holds two messages (x_0, x_1) , and a *receiver* who wishes to learn exactly one of these messages x_b based on their choice bit $b \in \{0, 1\}$. The protocol must satisfy two crucial security properties:

- **Receiver Privacy:** The sender should remain oblivious to which message was chosen (i.e., the value of b).
- **Sender Privacy:** The receiver should learn only one message x_b and gain no information about the other message x_{1-b} .

1.2 Applications and Significance

The importance of OT extends far beyond its simple description, playing a crucial role in:

- **Secure Multi-party Computation (MPC):** OT serves as the foundation for general secure computation protocols, enabling multiple parties to jointly compute functions while keeping their inputs private.
- **Private Set Intersection (PSI):** Protocols allowing two parties to find common elements in their sets without revealing other elements heavily rely on OT as a building block.

- **Zero-Knowledge Proofs:** Many efficient zero-knowledge protocols utilize OT as a fundamental component to achieve their security properties.
- **Password-Based Authentication:** OT can be used to construct secure password authentication protocols that protect against offline dictionary attacks.

1.3 Historical Development

The evolution of OT protocols represents a fascinating journey in cryptographic research:

- **1981:** Rabin introduced the original concept of OT where the receiver receives a message with probability $1/2$.
- **1985:** Even, Goldreich, and Lempel proposed the 1-out-of-2 variant, which proved more practical for applications.
- **1992:** Bellare and Micali presented their DDH-based protocol, offering a simple and elegant solution.
- **Recent Developments:** Modern implementations focus on efficiency improvements through techniques like OT extension and batch processing.

The Bellare-Micali protocol [1] represents a particularly elegant implementation of 1-out-of-2 OT. Its security relies on the well-studied Decisional Diffie-Hellman (DDH) assumption in cyclic groups, making it both theoretically sound and practically implementable. The protocol cleverly combines ElGamal-style encryption with hashing techniques to achieve its security goals.

1.4 Notation and Assumptions

To proceed with the formal description of the protocol, we establish the following notation and assumptions:

- Let p be a large prime, typically chosen to provide security equivalent to standard cryptographic parameters (e.g., 2048 bits for classical security).
- Let G be a cyclic group of order p with a generator g . Common instantiations include:
 - A prime-order subgroup of \mathbb{Z}_p^*
 - An elliptic curve group over a finite field
 - Other cryptographic groups where DDH is believed to be hard
- We denote by $H(\cdot)$ a cryptographic hash function that maps elements of G (or bitstrings derived thereof) to $\{0,1\}^n$ for some security parameter n . This function is modeled as a random oracle in security proofs.
- We assume the DDH assumption holds in G . Formally, for random $a, b \in \mathbb{Z}_p$, given (g, g^a, g^b) , it is computationally infeasible to distinguish g^{ab} from a random element in G .
- The sender, denoted as A , has two messages: $x_0, x_1 \in \{0,1\}^n$.
- The receiver, denoted as B , has a choice bit $b \in \{0,1\}$ indicating which message they want to learn.

1.5 Security Goals

The protocol aims to achieve several security properties:

- **Correctness:** If both parties follow the protocol, the receiver should successfully obtain their chosen message with probability 1.
- **Receiver Privacy:** The sender should learn nothing about the receiver's choice bit b , even if they deviate from the protocol (malicious behavior).
- **Sender Privacy:** The receiver should learn exactly one message and nothing about the other message, even under malicious behavior.
- **Simulatability:** The protocol should be simulatable in both the semi-honest and malicious security models.

2 Protocol Definition

2.1 High-Level Idea

At a high level, the Bellare-Micali protocol implements oblivious transfer through a clever application of public key cryptography. The core insight lies in creating a pair of public keys with special properties:

- The receiver generates two public keys (PK_0, PK_1) but can only know the secret key for one of them
- The public keys are mathematically related such that their product equals a public value
- To the sender, both public keys appear equally valid and indistinguishable

The protocol achieves this through several elegant steps:

1. The sender first publishes a random group element $c = g^u$
2. For choice bit b , the receiver creates:
 - A "real" public key $PK_b = g^k$ for which they know the secret key k
 - A "fake" public key $PK_{1-b} = c/g^k$ for which they cannot know the secret key
3. This construction ensures $PK_0 \cdot PK_1 = c$, making both keys look valid to the sender
4. The DDH assumption ensures the sender cannot tell which key is "real" and which is "fake"

2.2 Protocol Steps

1. Setup:

- The sender A publishes a random element $c \in G$.

$$c = g^u \quad \text{for some random } u \in \mathbb{Z}_p.$$

2. Receiver's Key Generation:

- B chooses a random exponent $k \in \mathbb{Z}_p$.
- B computes

$$PK_b = g^k, \quad PK_{1-b} = \frac{c}{g^k} = g^u \cdot g^{-k} = g^{u-k}.$$

- Note that $PK_0 \cdot PK_1 = g^k \cdot g^{u-k} = g^u = c$, verifying the multiplication check.
- B sends (PK_0, PK_1) to A .
- A checks that $PK_0 \cdot PK_1 = c$. If the equality fails, A aborts.

3. Sender's Encryption:

- A samples random exponents $r_0, r_1 \in \mathbb{Z}_p$.
- A then computes two ciphertexts (C_0, C_1) :

$$C_0 = (g^{r_0}, H((PK_0)^{r_0}) \oplus x_0), \quad C_1 = (g^{r_1}, H((PK_1)^{r_1}) \oplus x_1).$$

- A sends (C_0, C_1) to B .

4. Receiver's Decryption:

- Upon receiving (C_0, C_1) , B decrypts the ciphertext corresponding to its choice bit b :

$$C_b = (V_1, V_2) = (g^{r_b}, H((PK_b)^{r_b}) \oplus x_b).$$

- Since $PK_b = g^k$, we have $(PK_b)^{r_b} = (g^k)^{r_b} = g^{kr_b}$.
- Let $V_1 = g^{r_b}$. Then,

$$V_1^k = (g^{r_b})^k = g^{kr_b}.$$

- B recovers

$$x_b = H(g^{kr_b}) \oplus (H(g^{kr_b}) \oplus x_b).$$

- This simplifies to x_b , as the XOR of y with itself is 0: $y \oplus y = 0$.

2.3 Pseudocode

Algorithm 1 Bellare-Micali 1-out-of-2 Oblivious Transfer

1: Input:

- Sender A has messages $x_0, x_1 \in \{0, 1\}^n$.
- Receiver B has choice bit $b \in \{0, 1\}$.

2: Output:

- B learns x_b .
- A remains oblivious to b .

3: Setup: A samples a random $u \in \mathbb{Z}_p$ and sets $c = g^u$. A sends c to B .

4: Receiver's KeyGen:

- B picks a random $k \in \mathbb{Z}_p$.
- B computes $PK_b = g^k$ and $PK_{1-b} = c / g^k$.
- B sends (PK_0, PK_1) to A .

5: Check: A verifies that $PK_0 \cdot PK_1 = c$. If it holds, proceed; otherwise, abort.

6: Sender's Encryption:

- A picks random exponents $r_0, r_1 \in \mathbb{Z}_p$.
- A computes:

$$C_0 \leftarrow \left(g^{r_0}, H((PK_0)^{r_0}) \oplus x_0 \right),$$

$$C_1 \leftarrow \left(g^{r_1}, H((PK_1)^{r_1}) \oplus x_1 \right).$$

- A sends (C_0, C_1) to B .

7: Receiver's Decryption:

- B takes $C_b = (V_1, V_2)$.
- B computes $V_1^k = (g^{r_b})^k = g^{kr_b}$.
- B recovers $x_b \leftarrow H(g^{kr_b}) \oplus V_2$.

8: Output: B obtains x_b .

3 Correctness Proof

Theorem 3.1: Correctness

The Bellare-Micali protocol satisfies the following correctness properties:

1. The receiver B successfully learns x_b for its chosen bit b

2. The protocol reveals no additional information about x_{1-b}
3. The decryption process is guaranteed to work when both parties follow the protocol

Proof of Correctness. We prove correctness by analyzing each component of the protocol systematically:

Part 1: Validity of Public Key Construction

First, we verify that the public keys are properly constructed:

- For the chosen bit b , B sets $PK_b = g^k$ where k is known
- The other key is set as $PK_{1-b} = c/g^k = g^u/g^k = g^{u-k}$
- This ensures $PK_0 \cdot PK_1 = g^k \cdot g^{u-k} = g^u = c$

Part 2: Analysis of Encryption

The sender creates two ciphertexts:

$$C_0 = (g^{r_0}, H((PK_0)^{r_0}) \oplus x_0) \quad \text{and} \quad C_1 = (g^{r_1}, H((PK_1)^{r_1}) \oplus x_1)$$

Let's examine why this construction works:

1. For the chosen bit b :

$$\begin{aligned} C_b &= (V_1, V_2) \\ &= (g^{r_b}, H((PK_b)^{r_b}) \oplus x_b) \\ &= (g^{r_b}, H((g^k)^{r_b}) \oplus x_b) \\ &= (g^{r_b}, H(g^{kr_b}) \oplus x_b) \end{aligned}$$

2. The receiver can compute:

$$\begin{aligned} V_1^k &= (g^{r_b})^k \\ &= g^{kr_b} \end{aligned}$$

Part 3: Decryption Process

The decryption succeeds because:

$$H((g^{r_b})^k) = H(g^{kr_b}) = H((PK_b)^{r_b})$$

Therefore:

$$\begin{aligned} x_b &= \left(H(g^{kr_b}) \oplus x_b \right) \oplus H(g^{kr_b}) \\ &= x_b \oplus \underbrace{H(g^{kr_b}) \oplus H(g^{kr_b})}_{=0} \\ &= x_b \end{aligned}$$

Part 4: Why No Information about x_{1-b} is Leaked

For the unchosen bit $(1 - b)$, the receiver cannot decrypt C_{1-b} because:

- They don't know the discrete logarithm of PK_{1-b} with respect to g
- Due to the DDH assumption, $(PK_{1-b})^{r_{1-b}}$ appears random
- Therefore $H((PK_{1-b})^{r_{1-b}})$ acts as a one-time pad
- This makes x_{1-b} information-theoretically secure

Thus, we have shown that:

1. The receiver can successfully decrypt their chosen message x_b
2. The decryption process is efficient and deterministic
3. No information about x_{1-b} is revealed
4. The protocol achieves perfect correctness when followed honestly

□

4 Security Analysis

In this section, we provide an informal overview of security based on the standard assumptions. A fully rigorous proof would require reductions to the DDH or CDH problem in the random oracle model (depending on the exact variant and threat model).

4.1 Sender's Security

Goal: Show that a (possibly cheating) receiver B cannot learn both x_0 and x_1 with non-negligible probability under the DDH (or CDH) assumption.

Theorem 4.1: Sender's Security

Assume that the Decisional Diffie-Hellman (DDH) problem is hard in the group G . Then, in the Bellare-Micali 1-out-of-2 OT protocol, no polynomial-time receiver B can learn both messages x_0 and x_1 from the ciphertexts (C_0, C_1) with non-negligible probability.

Proof. We sketch a standard reduction-based argument. Suppose, for contradiction, that there exists a probabilistic polynomial-time (PPT) adversary B^* (a malicious receiver) that learns *both* x_0 and x_1 with non-negligible advantage. We show how to construct a PPT algorithm \mathcal{D} that solves the DDH problem in G with non-negligible probability, contradicting the DDH assumption.

DDH Setup. Recall that the DDH problem in G is: given a random tuple (g, g^a, g^b, T) , decide whether $T = g^{ab}$ or T is a random element in G . We assume \mathcal{D} is given an input tuple (g, g^a, g^b, T) and must determine if $T = g^{ab}$. We will use this tuple to set up the OT protocol so that if B^* can learn both x_0 and x_1 , then \mathcal{D} can distinguish g^{ab} from a random element.

Simulation. The reduction \mathcal{D} acts as the sender A in the OT protocol, interacting with B^* .

1. **Setup:** \mathcal{D} chooses a random exponent $u \in \mathbb{Z}_p$ and sets

$$c = g^u \quad (\text{the published group element}).$$

\mathcal{D} sends c to B^* .

2. **Key Generation by B^* :** The adversarial B^* returns (PK_0, PK_1) such that $PK_0 \cdot PK_1 = c$ (or it may deviate arbitrarily if it tries to cheat).
3. **Encryption Phase:** The reduction \mathcal{D} must produce the ciphertexts C_0 and C_1 . Typically, A would choose $r_0, r_1 \in \mathbb{Z}_p$ uniformly at random and compute

$$C_0 = (g^{r_0}, H((PK_0)^{r_0}) \oplus x_0), \quad C_1 = (g^{r_1}, H((PK_1)^{r_1}) \oplus x_1).$$

Here is where \mathcal{D} embeds the DDH challenge into one of these components, ensuring that if B^* decrypts both, \mathcal{D} learns something about T being g^{ab} or random.

Concretely, \mathcal{D} can, for instance, set:

$$g^{r_0} \leftarrow g^a, \quad g^{r_1} \leftarrow g^b, \quad \text{and} \quad H((PK_0)^{r_0}) \leftarrow H((PK_0)^a), \quad H((PK_1)^{r_1}) \leftarrow H((PK_1)^b).$$

If $PK_0 = g^k$, then $(PK_0)^a = g^{ak}$; similarly, if $PK_1 = g^{u-k}$, then $(PK_1)^b = g^{b(u-k)}$, etc. The key point is that \mathcal{D} can craft the XORed part to embed T (from the DDH challenge) in one of the two ciphertexts. For example, \mathcal{D} might define

$$H((PK_0)^a) \oplus x_0 = \text{some function of } T, \quad H((PK_1)^b) \oplus x_1 = \text{some function of } T.$$

4. **Adversarial Decryption:** If B^* truly can obtain both x_0 and x_1 , then B^* effectively learns enough information about $(PK_0)^a$ and $(PK_1)^b$ to distinguish whether $T = g^{ab}$ or a random element. Specifically, knowledge of both x_0 and x_1 reveals which hashed values correspond to valid Diffie-Hellman components.

Extracting the DDH Decision. If $T = g^{ab}$, then the embedded encryption values will yield consistent decryptions that B^* can exploit to recover both x_0 and x_1 . If T is random, there is a non-negligible chance that B^* 's attempt to recover both messages will fail (or produce inconsistent values) because the underlying exponentiations will not match any legitimate Diffie-Hellman product. By observing B^* 's success or failure, \mathcal{D} distinguishes between $T = g^{ab}$ and T random, thereby solving the DDH problem with non-negligible probability.

Conclusion. Hence, if B^* succeeds in learning both messages with non-negligible probability, then \mathcal{D} breaks the DDH assumption. Since DDH is assumed to be hard in G , no such B^* can exist. Therefore, the Bellare-Micali OT protocol is secure against a malicious receiver trying to learn both messages. \square

4.2 Receiver's Security

Goal: Show that the sender A cannot determine which bit $b \in \{0, 1\}$ is chosen by B .

Theorem 4.2: Receiver's Security

Under the Decisional Diffie-Hellman assumption, a sender A (even if malicious) cannot distinguish which of the two public keys PK_0 or PK_1 is associated with the receiver's actual choice bit b . Consequently, A cannot learn b with non-negligible advantage.

Proof. We again use a standard reduction argument. Suppose there exists a PPT sender A^* that can determine the choice bit b of the receiver B with non-negligible advantage. We construct a PPT algorithm \mathcal{D} that breaks the DDH assumption.

DDH Setup. Let (g, g^x, g^y, T) be an instance of the DDH problem in G . The goal is to decide whether $T = g^{xy}$ or T is uniformly random in G .

Simulation as the Receiver. The algorithm \mathcal{D} will simulate B 's role and interact with A^* (the alleged distinguisher of b):

1. **Publishing c :** Typically, the sender A picks a random $u \in \mathbb{Z}_p$ and sends $c = g^u$ to B . In our simulation, \mathcal{D} does nothing special here; it just obtains c from A^* . (Alternatively, \mathcal{D} might control c if needed, but we can adapt how we embed the challenge.)
2. **Receiver Key Generation:** In the real protocol, the receiver chooses $k \in \mathbb{Z}_p$ and sets

$$PK_b = g^k, \quad PK_{1-b} = c/g^k.$$

The key point for security is that (PK_0, PK_1) is a random pair of group elements constrained by the product $PK_0 \cdot PK_1 = c$. Under the DDH assumption, to an outside observer (including A), (PK_0, PK_1) appears random (subject to the product constraint).

To embed the DDH challenge, \mathcal{D} can decide which one of g^x or g^y stands for PK_0 or PK_1 , effectively randomizing the assignment so that if A^* can guess b , it must be solving a DDH instance.

3. **Interaction and Challenge:** The malicious sender A^* attempts to determine b . But from A^* 's perspective, (PK_0, PK_1) is just some pair of elements in G such that $PK_0 \cdot PK_1 = c$. If \mathcal{D} picks $PK_b = g^x$ and $PK_{1-b} = c/g^x$, the distribution of these values is indistinguishable from the real protocol if T is indeed g^{xy} . If T is random, we can adjust the scenario (for instance, one of these might not be a correct Diffie-Hellman share). In either case, any advantage A^* has in guessing b would yield an advantage in distinguishing $T = g^{xy}$ from T random.

Using A^* 's Advantage to Break DDH. If A^* truly can distinguish b with non-negligible probability, then by controlling how we assign (PK_0, PK_1) in terms of (g^x, g^y) and $c = g^u$, we can make A^* 's decision reveal information about whether $T = g^{xy}$ or random. Specifically:

- If T is a valid Diffie-Hellman term ($T = g^{xy}$), then the distribution of (PK_0, PK_1) is consistent with one fixed b .
- If T is random, then from A^* 's view, the assignment is inconsistent with a real Diffie-Hellman triple, changing the probability that A^* guesses b correctly.

Hence, any non-negligible bias in A^* 's guess of b translates into a non-negligible advantage in deciding whether $T = g^{xy}$, thus solving the DDH problem.

Conclusion. This contradiction implies that no polynomial-time sender A^* can guess the receiver's bit b with advantage better than $1/2 + \text{negl}(p)$. Hence, the protocol hides b from A , establishing receiver's security under the DDH assumption. \square

4.3 Extensions and the Malicious Setting

A subtle point arises when the receiver is *malicious* (i.e., not following the protocol exactly). In such a scenario, the receiver might attempt to craft (PK_0, PK_1) in such a way to gain partial information about both x_0 and x_1 .

Remark 4.3

When modeling H as a random oracle, one can strengthen the security argument to show that *any* adversarial strategy that breaks the OT properties would imply the ability to solve the Computational Diffie-Hellman (CDH) problem. For a more thorough treatment of malicious security, see [3].

5 Conclusion

The Bellare-Micali 1-out-of-2 Oblivious Transfer protocol demonstrates how classical public key techniques (like ElGamal) combined with hashing can yield a simple, elegant solution for oblivious transfer under the Decisional Diffie-Hellman assumption.

Key Points:

- The receiver's *choice bit* b is hidden due to the indistinguishability of PK_b vs. PK_{1-b} under DDH.
- The sender's messages x_0, x_1 remain protected, as the receiver can only decrypt one ciphertext properly.
- Extensions exist to handle more robust (e.g., malicious) adversaries in the random oracle model.

Open Problems / Future Directions:

- Adapting this protocol to post-quantum settings where DDH may be weaker.
- Extending to protocols that require *multiple* oblivious transfers or more complex secure computations.
- Practical optimizations, such as using elliptic curves or tailored hash functions for performance gains.

References

-
- [1] Bellare, M., and Micali, S. “Oblivious Transfer and Polynomial Evaluation.” *Advances in Cryptology—CRYPTO ’92*, Springer, 1992.
 - [2] Diffie, W., and Hellman, M. “New Directions in Cryptography.” *IEEE Transactions on Information Theory*, 1976.
 - [3] Naor, M., and Pinkas, B. “Oblivious Transfer and Polynomial Evaluation.” *Proceedings of the Thirty-First Annual ACM Symposium on Theory of Computing (STOC)*, 1999.

A Mathematical Details and Implementation Considerations

A.1 Properties of XOR Operation

The protocol’s use of XOR (\oplus) is crucial for several reasons:

- **Perfect Secrecy:** When used with a random key of equal length, XOR provides information-theoretic security:

$$\forall m, c : \Pr[m \oplus k = c] = 2^{-n}$$

- **Computational Efficiency:** XOR operations are extremely fast in hardware:

- Single CPU cycle per word-size operation
- Parallelizable across multiple bits
- No carry propagation needed

- **Self-Inverse Property:** XOR simplifies decryption because:

$$(m \oplus k) \oplus k = m$$

A.2 Group Selection and Security Implications

The choice of group G significantly impacts security and performance:

Group Type	Advantages	Disadvantages	Security Level
\mathbb{Z}_p^* subgroups	- Well understood - Simple implementation	- Large parameters - Slower operations	3072-bit \approx 128-bit security
Elliptic Curves	- Smaller parameters - Faster operations	- Complex implementation - Point validation needed	256-bit \approx 128-bit security

Table 1: Comparison of Common Group Choices

A.3 Parameter Size Recommendations

Security levels should be chosen based on the intended application:

- **Classical Security Levels:**

80-bit: $p \approx 1024$ bits or 160-bit EC
128-bit: $p \approx 3072$ bits or 256-bit EC
256-bit: $p \approx 15360$ bits or 512-bit EC

- **Post-Quantum Considerations:** Parameters should be doubled for quantum resistance

A.4 Efficiency Analysis

Cost breakdown per operation:

- **Exponentiation:** $O(\log p)$ multiplications
- **Hash Function:** $O(1)$ for fixed input size
- **XOR Operation:** $O(n)$ bit operations for n -bit messages

Total computational complexity:

$$O(\log p) \text{ multiplications} + O(n) \text{ bit operations}$$

A.5 Implementation Considerations

Critical aspects for secure implementation:

1. **Random Number Generation:**

- Use cryptographically secure RNG
- Fresh randomness for each protocol run
- Proper entropy collection

2. **Hash Function Requirements:**

- Collision resistance: $2^{n/2}$ security
- Pre-image resistance: 2^n security
- Output length matching message size

3. **Group Element Validation:**

- Verify elements are in the correct subgroup
- Check order requirements
- Reject invalid inputs

A.6 Performance Optimizations

Common techniques for improving efficiency:

- **Pre-computation:**
 - Fixed-base exponentiation tables
 - Window methods for exponentiation
 - Batch processing for multiple transfers
- **Implementation Tricks:**
 - Montgomery multiplication
 - Karatsuba multiplication for large numbers
 - Simultaneous multiple exponentiation

B Security Analysis Extensions

B.1 Malicious Adversary Considerations

Additional security measures for malicious adversaries:

- **Zero-Knowledge Proofs:**
 - Proof of knowledge for discrete logarithms
 - Proof of proper key generation
 - Range proofs for parameters
- **Protocol Augmentation:**
 - Commitment schemes
 - Cut-and-choose techniques
 - Verifiable random functions

B.2 Advanced Security Properties

Additional security features achievable with modifications:

- **Forward Secrecy:**
 - Ephemeral keys
 - Session key derivation
 - Perfect forward secrecy options
- **Fairness:**
 - Gradual release mechanisms
 - Timeout-based aborts
 - Trusted setup options